

Simpson's paradox

Data Science in a Box
datasciencebox.org



Case study: Berkeley admission data



Berkeley admission data

- Study carried out by the Graduate Division of the University of California, Berkeley in the early 70's to evaluate whether there was a gender bias in graduate admissions.
- The data come from six departments. For confidentiality we'll call them A-F.
- We have information on whether the applicant was male or female and whether they were admitted or rejected.
- First, we will evaluate whether the percentage of males admitted is indeed higher than females, overall. Next, we will calculate the same percentage for each department.



Data

```
## # A tibble: 4,526 × 3
##   admit  gender dept
##   <fct>  <fct> <ord>
## 1 Admitted Male   A
## 2 Admitted Male   A
## 3 Admitted Male   A
## 4 Admitted Male   A
## 5 Admitted Male   A
## 6 Admitted Male   A
## 7 Admitted Male   A
## 8 Admitted Male   A
## 9 Admitted Male   A
## 10 Admitted Male  A
## 11 Admitted Male  A
## 12 Admitted Male  A
## 13 Admitted Male  A
## 14 Admitted Male  A
## 15 Admitted Male  A
## # ... with 4,511 more rows
```

```
## # A tibble: 2 × 2
##   gender    n
##   <fct> <int>
## 1 Female  1835
## 2 Male    2691
```

```
## # A tibble: 6 × 2
##   dept    n
##   <ord> <int>
## 1 A      933
## 2 B      585
## 3 C      918
## 4 D      792
## 5 E      584
## 6 F      714
```

```
## # A tibble: 2 × 2
##   admit    n
##   <fct> <int>
## 1 Rejected 2771
## 2 Admitted 1755
```



What can you say about the overall gender distribution? Hint: Calculate the following probabilities: $P(\text{Admit}|\text{Male})$ and $P(\text{Admit}|\text{Female})$.

```
ucbadmit %>%  
  count(gender, admit)
```

```
## # A tibble: 4 × 3  
##   gender admit      n  
##   <fct> <fct>   <int>  
## 1 Female Rejected  1278  
## 2 Female Admitted   557  
## 3 Male   Rejected  1493  
## 4 Male   Admitted  1198
```



```
ucbadmit %>%  
  count(gender, admit) %>%  
  group_by(gender) %>%  
  mutate(prop_admit = n / sum(n))
```

```
## # A tibble: 4 × 4  
## # Groups:   gender [2]  
##   gender admit      n prop_admit  
##   <fct> <fct>   <int>      <dbl>  
## 1 Female Rejected  1278      0.696  
## 2 Female Admitted   557      0.304  
## 3 Male   Rejected  1493      0.555  
## 4 Male   Admitted  1198      0.445
```


- $P(\text{Admit}|\text{Female}) = 0.304$
- $P(\text{Admit}|\text{Male}) = 0.445$



Overall gender distribution

Plot

Code



Overall gender distribution

Plot

Code

```
ggplot(ucbadmit, aes(y = gender, fill = admit)) +  
  geom_bar(position = "fill") +  
  labs(title = "Admit by gender",  
        y = NULL, x = NULL)
```



What can you say about the gender distribution by department ?

```
ucbadmit %>%  
  count(dept, gender, admit)
```

```
## # A tibble: 24 × 4  
##   dept gender admit      n  
##   <ord> <fct>  <fct>   <int>  
## 1 A      Female Rejected    19  
## 2 A      Female Admitted    89  
## 3 A      Male   Rejected   313  
## 4 A      Male   Admitted   512  
## 5 B      Female Rejected     8  
## 6 B      Female Admitted    17  
## # ... with 18 more rows
```



Let's try again... What can you say about the gender distribution by department?

```
ucbadmit %>%  
  count(dept, gender, admit) %>%  
  pivot_wider(names_from = dept, values_from = n)
```


```
## # A tibble: 4 × 8  
##   gender admit      A      B      C      D      E      F  
##   <fct> <fct>   <int> <int> <int> <int> <int> <int>  
## 1 Female Rejected    19      8   391   244   299   317  
## 2 Female Admitted    89     17   202   131    94    24  
## 3 Male   Rejected   313   207   205   279   138   351  
## 4 Male   Admitted   512   353   120   138    53    22
```



Gender distribution, by department

Plot

Code



Gender distribution, by department



Plot

Code

```
ggplot(ucbadmit, aes(y = gender, fill = admit)) +  
  geom_bar(position = "fill") +  
  facet_wrap(. ~ dept) +  
  scale_x_continuous(labels = label_percent()) +  
  labs(title = "Admissions by gender and department",  
       x = NULL, y = NULL, fill = NULL) +  
  theme(legend.position = "bottom")
```



Case for gender discrimination?



Closer look at departments

Output

Code

```
## # A tibble: 12 × 5
## # Groups:   dept, gender [12]
##   dept gender n_admitted n_applied prop_admit
##   <ord> <fct>      <int>      <int>      <dbl>
## 1 A     Female      89       108      0.824
## 2 A     Male      512      825      0.621
## 3 B     Female      17        25      0.68
## 4 B     Male     353      560      0.630
## 5 C     Female     202      593      0.341
## 6 C     Male     120      325      0.369
## 7 D     Female     131      375      0.349
## 8 D     Male     138      417      0.331
## 9 E     Female      94      393      0.239
## 10 E    Male       53      191      0.277
## 11 F    Female      24      341      0.0704
## 12 F    Male       22      373      0.0590
```



Closer look at departments

Output

Code

```
ucbadmit %>%  
  count(dept, gender, admit) %>%  
  group_by(dept, gender) %>%  
  mutate(  
    n_applied = sum(n),  
    prop_admit = n / n_applied  
  ) %>%  
  filter(admit == "Admitted") %>%  
  rename(n_admitted = n) %>%  
  select(-admit) %>%  
  print(n = 12)
```




Simpson's paradox




Relationship between two variables

```
## # A tibble: 8 × 3
##       x     y z
##   <dbl> <dbl> <chr>
## 1     2     4 A
## 2     3     3 A
## 3     4     2 A
## 4     5     1 A
## 5     6    11 B
## 6     7    10 B
## # ... with 2 more rows
```




Relationship between two variables

```
## # A tibble: 8 × 3
##       x     y z
##   <dbl> <dbl> <chr>
## 1     2     4 A
## 2     3     3 A
## 3     4     2 A
## 4     5     1 A
## 5     6    11 B
## 6     7    10 B
## # ... with 2 more rows
```




Considering a third variable

```
## # A tibble: 8 × 3
##       x     y z
##   <dbl> <dbl> <chr>
## 1     2     4 A
## 2     3     3 A
## 3     4     2 A
## 4     5     1 A
## 5     6    11 B
## 6     7    10 B
## # ... with 2 more rows
```



Relationship between three variables

```
## # A tibble: 8 × 3
##       x     y z
##   <dbl> <dbl> <chr>
## 1     2     4 A
## 2     3     3 A
## 3     4     2 A
## 4     5     1 A
## 5     6    11 B
## 6     7    10 B
## # ... with 2 more rows
```



Simpson's paradox

- Not considering an important variable when studying a relationship can result in **Simpson's paradox**
- Simpson's paradox illustrates the effect that omission of an explanatory variable can have on the measure of association between another explanatory variable and a response variable
- The inclusion of a third variable in the analysis can change the apparent relationship between the other two variables



Aside: `group_by()` and `count()`



What does group_by() do?

`group_by()` takes an existing data frame and converts it into a grouped data frame where subsequent operations are performed "once per group"

```
ucbadmit
```

```
## # A tibble: 4,526 × 3
##   admit    gender dept
##   <fct>    <fct> <ord>
## 1 Admitted Male    A
## 2 Admitted Male    A
## 3 Admitted Male    A
## 4 Admitted Male    A
## 5 Admitted Male    A
## 6 Admitted Male    A
## # ... with 4,520 more rows
```

```
ucbadmit %>%
  group_by(gender)
```

```
## # A tibble: 4,526 × 3
## # Groups:   gender [2]
##   admit    gender dept
##   <fct>    <fct> <ord>
## 1 Admitted Male    A
## 2 Admitted Male    A
## 3 Admitted Male    A
## 4 Admitted Male    A
## 5 Admitted Male    A
## 6 Admitted Male    A
## # ... with 4,520 more rows
```



What does group_by() not do?

`group_by()` does not sort the data, `arrange()` does

```
ucbadmit %>%  
  group_by(gender)
```

```
## # A tibble: 4,526 × 3  
## # Groups:   gender [2]  
##   admit    gender dept  
##   <fct>    <fct> <ord>  
## 1 Admitted Male   A  
## 2 Admitted Male   A  
## 3 Admitted Male   A  
## 4 Admitted Male   A  
## 5 Admitted Male   A  
## 6 Admitted Male   A  
## # ... with 4,520 more rows
```

```
ucbadmit %>%  
  arrange(gender)
```

```
## # A tibble: 4,526 × 3  
##   admit    gender dept  
##   <fct>    <fct> <ord>  
## 1 Admitted Female A  
## 2 Admitted Female A  
## 3 Admitted Female A  
## 4 Admitted Female A  
## 5 Admitted Female A  
## 6 Admitted Female A  
## # ... with 4,520 more rows
```



What does group_by() not do?

`group_by()` does not create frequency tables, `count()` does

```
ucbadmit %>%  
  group_by(gender)
```

```
## # A tibble: 4,526 × 3  
## # Groups:   gender [2]  
##   admit    gender dept  
##   <fct>    <fct> <ord>  
## 1 Admitted Male   A  
## 2 Admitted Male   A  
## 3 Admitted Male   A  
## 4 Admitted Male   A  
## 5 Admitted Male   A  
## 6 Admitted Male   A  
## # ... with 4,520 more rows
```

```
ucbadmit %>%  
  count(gender)
```

```
## # A tibble: 2 × 2  
##   gender      n  
##   <fct>  <int>  
## 1 Female  1835  
## 2 Male   2691
```



Undo grouping with ungroup()

```
ucbadmit %>%  
  count(gender, admit) %>%  
  group_by(gender) %>%  
  mutate(prop_admit = n / sum(n)) %>%  
  select(gender, prop_admit)
```

```
## # A tibble: 4 × 2  
## # Groups:   gender [2]  
##   gender prop_admit  
##   <fct>      <dbl>  
## 1 Female      0.696  
## 2 Female      0.304  
## 3 Male        0.555  
## 4 Male        0.445
```

```
ucbadmit %>%  
  count(gender, admit) %>%  
  group_by(gender) %>%  
  mutate(prop_admit = n / sum(n)) %>%  
  select(gender, prop_admit) %>%  
  ungroup()
```

```
## # A tibble: 4 × 2  
##   gender prop_admit  
##   <fct>      <dbl>  
## 1 Female      0.696  
## 2 Female      0.304  
## 3 Male        0.555  
## 4 Male        0.445
```



count() is a short-hand

`count()` is a short-hand for `group_by()` and then `summarise()` to count the number of observations in each group

```
ucbadmit %>%  
  group_by(gender) %>%  
  summarise(n = n())
```

```
## # A tibble: 2 × 2  
##   gender      n  
##   <fct>   <int>  
## 1 Female  1835  
## 2 Male    2691
```

```
ucbadmit %>%  
  count(gender)
```

```
## # A tibble: 2 × 2  
##   gender      n  
##   <fct>   <int>  
## 1 Female  1835  
## 2 Male    2691
```



count can take multiple arguments

```
ucbadmit %>%  
  group_by(gender, admit) %>%  
  summarise(n = n())
```

```
## # A tibble: 4 × 3  
## # Groups:   gender [2]  
##   gender admit      n  
##   <fct> <fct>    <int>  
## 1 Female Rejected 1278  
## 2 Female Admitted  557  
## 3 Male   Rejected 1493  
## 4 Male   Admitted 1198
```

```
ucbadmit %>%  
  count(gender, admit)
```

```
## # A tibble: 4 × 3  
##   gender admit      n  
##   <fct> <fct>    <int>  
## 1 Female Rejected 1278  
## 2 Female Admitted  557  
## 3 Male   Rejected 1493  
## 4 Male   Admitted 1198
```



summarise() after group_by()

- count() ungroups after itself
- summarise() peels off one layer of grouping by default, or you can specify a different behaviour

```
ucbadmit %>%  
  group_by(gender, admit) %>%  
  summarise(n = n())
```

```
## # A tibble: 4 × 3  
## # Groups:   gender [2]  
##   gender admit      n  
##   <fct>   <fct>   <int>  
## 1 Female Rejected  1278  
## 2 Female Admitted   557  
## 3 Male   Rejected  1493  
## 4 Male   Admitted  1198
```

